



Task 1: Mountains (mountains)

Authored and prepared by: Lim An Jun

Subtask 1

Limits: H_i is non-decreasing.

Since H_i is non-decreasing, there does not exist mountains y and z such that $y < z$ and $H_y > H_z$. Thus, there are no valid triplets and we output 0.

Subtask 2

Limits: $0 \leq H_i \leq 1$

All heights are 0 or 1, so $H_x = H_z = 0$ and $H_y = 1$. Let $smallLeft_i$ be the number of '0's in H_1, H_2, \dots, H_{i-1} and $smallRight_i$ be the number of '0's in $H_{i+1}, H_{i+2}, \dots, H_n$. Calculate and store their values in separate arrays in $O(n)$.

Consider a candidate for center mountain y where $H_y = 1$. There are $smallLeft_y$ candidates for x and $smallRight_y$ candidates for z . The number of choices is $smallLeft_y \times smallRight_y$. Calculate the sum of possible choices over all candidates y and output it.

Time complexity: $O(n)$

Subtask 3

Limits: $0 \leq H_i \leq 99$

Consider each possible height h for the center mountain y where $0 \leq h \leq 99$. Note that the magnitude of heights do not matter as we only care if each mountain is shorter than h or not. Treat all $H_i < h$ as 0 and all $H_i \geq h$ as 1. Then, the same approach to subtask 1 can be used.

Time complexity: $O(\max H_i \times n)$

Subtask 4

Limits: $0 \leq n \leq 500$

Iterate through all possible triplets of x, y, z such that $1 \leq x < y < z \leq n$ and count how many are valid, which takes $O(1)$ for each triplet.

Time complexity: $O(n^3)$



Subtask 5

Limits: $0 \leq n \leq 10^4$

If we copy the solution from subtask 3 directly, it will fail as $\max H_i = 10^{18}$. However, we only need to check $h = H_i$ for some $1 \leq i \leq n$. There are at most n such different values, and thus a modified solution to subtask 3 will pass.

Time complexity: $O(n^2)$

Subtask 6

Limits: $0 \leq H_i \leq 10^5$

So far, we have been re-counting $smallLeft_i$ and $smallRight_i$ naively for different heights. We can use a data structure that supports range queries and point updates with logarithmic complexity (e.g. fenwick tree, segment tree) to calculate them for all heights. Let $query(h)$ denote the number of heights that are at most h so far, and $update(h)$ adds a height h .

Iterate through H from left to right. At each index i , $smallLeft_i = query(H_i - 1)$. Then, we add in the current height with $update(H_i)$. $smallRight$ can be calculated similarly. The answer is the sum of $smallLeft_i \times smallRight_i$.

Time complexity: $O(n \log (\max H_i))$

Subtask 7

Limits: No additional constraints

The solution for subtask 6 fails as $\max H_i = 10^{18}$. However, note that only the relative orders of heights matter. Thus, we can discretise the heights into values within 1 to n (set smallest H_i to 1, second smallest to 2 etc.) and use the same approach.

Time complexity: $O(n \log n)$



Task 2: Visiting Singapore (visitingsingapore)

Authored by: Sung Wing Kin, Ken

Prepared by: Ng Yu Peng

Note: To avoid confusion, we will be treating A, B as positive integers throughout this writeup, so when we say a loss of A happiness we really mean $-A$ happiness algebraically.

Subtask 1

Limits: $K = 1, m \leq n \leq 10^3$

Since $K = 1$, the only event is event 1. As $m \leq n$, we can just attend the first m days all of which are event 1, and since we don't skip any events during the stay and we attend all events $T[1], T[2], \dots, T[m]$, no loss of happiness is incurred and our answer is simply $mV[1]$.

Time Complexity: $O(1)$

Subtask 2

Limits: $K = 1, n < m \leq 10^3$

Again the only event is 1. This time $n < m$ so we can only attend n events. So we will attend the first n days of events and miss $T[n+1], T[n+2], \dots, T[m]$. Hence the maximum happiness is $nV[1] - (n - m)B - A$.

Time Complexity: $O(1)$

Subtask 3

Limits: $A = B = 0$

For this subtask we will need to do some dynamic programming. We first treat each attended event as a pair of numbers: if we attend event $T[j]$ on day i , we denote it by the pair (i, j) . Let $dp(i, j)$ be the maximum happiness if the last event (x, y) we attend satisfies $x \leq i$ and $y \leq j$.

Since $A = B = 0$ we do not need to worry about incurring losses to happiness. If $S[i] = T[j]$ then we may choose to attend the event (i, j) . Otherwise we cannot attend (i, j) . Hence

$$dp(i, j) = \begin{cases} \max(dp(i-1, j-1) + V[T[j]], dp(i, j-1), dp(i-1, j)) & \text{if } S[i] = T[j] \\ \max(dp(i, j-1), dp(i-1, j)) & \text{if } S[i] \neq T[j] \end{cases}$$

Then our final answer is just $dp(n, m)$.

Time Complexity: $O(nm)$



Subtask 4

Limits: $A = 0$

We will still use dynamic programming, but now we need to change up the state. Redefine $dp(i, j)$ to represent the maximum happiness achievable if the last event (x, y) attended has $x \leq i$ and $y \leq j$, including the loss of happiness incurred from skipping events up to $S[i]$ and $T[j]$.

We first set $dp(i, j) = \max(dp(i, j-1) - B, dp(i-1, j) - B)$, to take into account the cases where we do not attend event (i, j) . In each case we additionally skip $T[j]$ and $S[i]$ respectively so we incur an additional loss of B .

Then if $S[i] = T[j]$, we have the option of attending the event (i, j) now. We can also have (i, j) be the first event we attend. Hence, if $f(i, j)$ is the maximum happiness achievable if the last event attended is (i, j) , ignoring costs from skipping events after $S[i]$ and $T[j]$, then

$$f(i, j) = \max(dp[i-1][j-1] + V[T[j]], V[T[j]] - (j-1)B)$$

Now we set if $f(i, j) > dp(i, j)$ then we set $dp(i, j) = f(i, j)$ as attending event (i, j) gives more happiness.

Now the answer is just the maximum of $f(i, j) - (m-j)B$ across all i, j , with the $-(m-j)B$ to account for the loss of happiness for skipping events after $T[j]$, as well as $-A - mB$ in the case we do not attend any events

Time Complexity: $O(nm)$

Subtask 5

Limits: $B = 0$

Now, if we attend event (i, j) , we need to know if we attended $S[i-1]$ and $T[j-1]$ as well to decide if we will incur an additional loss of A happiness. Let $dp(i, j, k)$ be the maximum happiness (not including losses for skipping events after $S[i], T[j]$) if the last event attended (x, y) has $x \leq i$ and $y \leq j$, where $k = 0$ has no restrictions, $k = 1$ means we attended $T[j]$, $k = 2$ means we attended $S[i]$, and $k = 3$ means we attended both $S[i]$ and $T[j]$.


We first set

$$dp(i, j, 0) = \max(dp(i, j-1, 0), dp(i-1, j, 0))$$

$$dp(i, j, 1) = dp(i-1, j, 1)$$

$$dp(i, j, 2) = dp(i, j-1, 2)$$

to note the cases where we do not attend event (i, j) .



Now if $S[i] = T[j]$, we need to consider the event (i, j) . We first consider the case where (i, j) is our first event. So firstly we set

$$dp(i, j, 3) = \begin{cases} V[T[j]] & \text{if } j = 1 \\ V[T[j]] - A & \text{if } j > 1 \end{cases}$$

Next we consider the case where there is a previous event. We incur a loss of A happiness for each of the events $S[i-1], T[j-1]$ we do not attend, so we consider the values $dp(i-1, j-1, 0) - 2*a + V[T[j]], dp(i-1, j-1, 1) - a + V[T[j]], dp(i-1, j-1, 2) - a + V[T[j]], dp(i-1, j-1, 3) + V[T[j]]$. If the maximum among these numbers is larger than the current value of $dp(i, j, 3)$, set $dp(i, j, 3)$ to this maximum. Hence now $dp(i, j, 3)$ is the maximum happiness attainable if we attend event (i, j) last, not considering penalties to happiness from events after $T[j]$.

With this we can now update our other dp values. If $dp(i, j, 3)$ is larger than any of the other $dp(i, j, k)$ for $k = 0, 1, 2$, set $dp(i, j, k) = dp(i, j, 3)$.

However, notice that as i, j can each go up to 5000, our memory usage will be around 400MB if we use int arrays. Hence we will need to optimise memory usage.

Note that to compute $dp(i, j, k)$ we only refer to indices $i-1, i, j-1, j$. Hence we can do DP on the fly and just use a $2 \times m \times 4$ array, with the first index representing $i-1$ and i in some order.

Now the answer is simply the maximum of $-A$ (don't attend any events), or the maximum among $dp(i, j, 3) - A$ for $j < m$ or the maximum among $dp(i, j, 3)$ for $j = m$.

Time Complexity: $O(nm)$

Subtask 6

Limits: $n, m < 100$

Let $f(i, j)$ be the maximum happiness achievable if the last event attended is (i, j) , ignoring costs from skipping events after $S[i]$ and $T[j]$, then we can just consider all the possible cases to compute $f(i, j)$. In particular, $f(i, j)$ is the maximum among all of

$$\begin{cases} f(i-1, j-1) + V[T[j]] & \\ f(i-1, j_1) + V[T[j]] - A - (j - j_1 - 1)B & \text{for } j_1 < j-1 \\ f(i_1, j-1) + V[T[j]] - A - (i - i_1 - 1)B & \text{for } i_1 < i-1 \\ f(i_1, j_1) + V[T[j]] - A - (i + j - i_1 - j_1 - 2)B & \text{for } j_1 < j-1 \text{ and } i_1 < i-1 \\ V[T[j]] & \text{if } j = 1 \\ V[T[j]] - A - (j-1)B & \text{if } j > 1 \end{cases}$$

Now the answer is just the maximum of $f(i, m)$ across all i , $f(i, j) - (m-j)B - A$ for all i, j with $j < m$, as well as $-A - mB$ in the case we do not attend any events.



Time Complexity: $O(n^2m^2)$

Subtask 7

Limits: No additional constraints

To solve the final subtask, we just need to merge our solutions for subtasks 4 and 5.

Everything is the same as in Subtask 5 except for the cases where we do not attend event (i, j) when considering $dp(i, j, k)$. So we first set

$$dp(i, j, 0) = \max(dp(i, j - 1, 0) - b, dp(i - 1, j, 0) - b)$$

$$dp(i, j, 1) = dp(i - 1, j, 1) - b$$

$$dp(i, j, 2) = dp(i, j - 1, 2) - b$$

to account for the additional loss of B happiness when we skip event $S[i]$ or $T[j]$. Then the rest of the solution follows as in Subtask 5.

Now the answer is simply the maximum of $-A - mB$ (don't attend any events), or the maximum among $dp(i, j, 3) - A - (m - j)B$ for $j < m$ or the maximum among $dp(i, j, 3)$ for $j = m$.

Time Complexity: $O(nm)$



Task 3: Solar Storm (**solarstorm**)

Authored and prepared by: Bernard Teo Zhi Yi

Subtask 1

Limits: $S = 1$, $N \leq 10^4$, $K \leq 10^9$, $d_i \leq 10^5$ for all $1 \leq i \leq N - 1$, $v_i \leq 10^5$ for all $1 \leq i \leq N$

There is exactly one shield. For each module, calculate the total value of all shielded modules if we deploy the shield there. This takes $O(N)$ per module, and we have N modules to check.

Time complexity: $O(N^2)$

Note: The constraints of this subtask are such that solutions which store cumulative distances and values using 32-bit signed integers will not overflow.

Subtask 2

Limits: $S = 1$, $d_i = 1$ for all $1 \leq i \leq N - 1$

There is exactly one shield, and furthermore the distance between any two adjacent modules is fixed at 1. This means that the optimal solution must protect a contiguous range of $2K + 1$ modules (unless $N \leq 2K$, in which case we can deploy the shield to protect all modules). We can solve this in $O(N)$ using the sliding window technique.

Time complexity: $O(N)$

Subtask 3

Limits: $S = 1$

This is similar to subtask 2, but now the distances between adjacent modules are no longer fixed.

For each module i , we can determine r_i , the rightmost module that will be protected if the shield is deployed at module i . This can be done in $O(N)$ with the sliding window technique. The modules can also be scanned in the reverse order to determine l_i , the leftmost module that will be protected if the shield is deployed at module i .

The optimal location to deploy the shield is then $\arg \max_{i \in \{1, \dots, N\}} \sum_{j=l_i}^{r_i} v_j$. The value $\sum_{j=l_i}^{r_i} v_j$ may be computed in $O(1)$ after computing the prefix sum of the v_i 's.

Time complexity: $O(N)$



Subtask 4

Limits: $K = 1$, $d_i = 2$ for all $1 \leq i \leq N - 1$

This is similar in essence to subtask 2. Each shield can only protect the module that it is deployed in. Any valid solution must hence deploy shields at S consecutive modules. We can check all possibilities in $O(N)$ using the sliding window technique.

Time complexity: $O(N)$

Subtask 5

Limits: $N \leq 10^4$

Given module i , let r_i be the rightmost module that is at most distance K away from module i (this is equivalent to the definition of r_i in subtask 3).

Suppose module i is the leftmost protected module. Then let $e_1(i) - 1$ be the rightmost module that is protected by the same shield that protects module i . $e_1(i) - 1$ can be determined by traversing the spaceship rightwards from module i . Then $e_1(i)$ is the leftmost module that is protected by the second shield (from the left), and we can similarly traverse the spaceship to find $e_2(i) - 1$, the rightmost module that is protected by the second shield. In this way, we can find $e_S(i) - 1$, the rightmost module that is protected by the S^{th} shield. At the same time, we keep track of the cumulative value of protected modules, so that after arriving at module $e_S(i) - 1$ we know the total value of protected modules. Since we traverse from left to right after starting from module i , it takes $O(N)$ to determine $e_S(i) - 1$ given a fixed i .

We try every i as a possible leftmost protected module – there are N modules that we could possibly start with.

Time complexity: $O(N^2)$

Note: Care should be taken to ensure that the implementation does not read past the right bound of the spaceship.

Subtask 6

Limits: $S \leq 10^2$

We can speed up the solution for subtask 5 by observing that we can compute the values of r_i once (using the sliding window technique described in subtask 3) and then cache these values, so that we can compute $e_S(i)$ in $O(S)$ time.

For each i , we can use the prefix sum method introduced in subtask 3 to compute the cumulative sum of values of modules in the range $[i, e_S(i))$.

Time complexity: $O(NS)$



Subtask 7

Limits: No additional constraints

We can build a directed graph of $N + 1$ vertices. Each vertex represents a module, and vertex $N + 1$ represents a dummy module past the right bound of the spaceship. For every vertex i ($1 \leq i \leq N$), add an edge from vertex i to vertex $e_1(i)$. Observe that the graph is a tree that is rooted at vertex $N + 1$ (i.e. all edges are directed towards vertex $N + 1$). Finding $e_S(i)$ is then equivalent to finding the S^{th} ancestor of the vertex i (if vertex i has depth less than S , then we define $e_S(i)$ to be vertex $N + 1$).

We then do a depth-first search on the tree starting from the root, while maintaining an auxiliary stack. When visiting vertex i , push i onto the stack, then visit its children, then pop i from the stack. Hence, when processing vertex i , the stack contains the list of vertices from i to the root (from top to bottom of the stack). The S^{th} element from the top of the stack is thus $e_S(i)$, which may be obtained in $O(1)$ if the stack is stored in a contiguous block of memory. Since the stack operations are $O(1)$, the whole depth-first search takes $O(N)$ time.

Lastly, we use the same prefix sum method as per subtask 3 and 6, which takes $O(1)$ time per range.

Time complexity: $O(N)$

Note: A solution that runs in $O(N \log N)$ is likely to get full points as well. There are a number of alternatives to the depth-first search which run in $O(N \log N)$, such as:

- Path compression on the tree. For each vertex i , we traverse S steps towards the root, and compress the path (i.e. replace the outgoing edge of each vertex j we traverse with an edge that goes from j directly to $e_S(i)$ (tagged with the number of steps the edge represents)). Since the function e_S is monotonic, we may process the vertices in increasing order of i (without fear of “skipping” the vertex that we want).
- Binary lifting on the tree. This is a well-known technique that can find the h^{th} ancestor of any given vertex in $O(\log h)$, after a preprocessing step that takes $O(N \log N)$ time.